

```
In [1]: import pyspark.pandas as ps
        from pyspark.sql import SparkSession
        import pandas as pd
        import numpy as np
        import warnings
        warnings.filterwarnings('ignore')
```

WARNING:root:'PYARROW_IGNORE_TIMEZONE' environment variable was not set. It is required to set this environment variable to '1' in both driver and executor sides if you use pyarrow>=2.0.0. pandas-on-Spark will set it for you but it does not work if there is a Spark context already launched.

Create Pandas on Spark DataFrame

```
In [2]: ps_df=ps.DataFrame([[ 'France', '50M', '3T'], [ 'India', '30M', '30T'], [ 'Kenya', '70M', '25T'],
                             [ 'Nigeria', '90M', '60T'], [ 'China', '20M', '2T'], [ 'USA', '80M', '30T'],
                             [ 'UK', '70M', '25T'], [ 'USA', '20M', '30T'], [ 'China', '70M', '25T'],
                             [ 'France', '50M', '3T'], [ 'China', '70M', '25T'] ],
                             columns=[ 'Country', 'Population', 'GDP'])

ps_df
```

Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties

Setting default log level to "WARN".

To adjust logging level use `sc.setLogLevel(newLevel)`. For SparkR, use `setLogLevel(newLevel)`.

22/01/14 11:47:59 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Out[2]:

	Country	Population	GDP
0	France	50M	3T
1	India	30M	30T
2	Kenya	70M	25T
3	Nigeria	90M	60T
4	China	20M	2T
5	USA	80M	30T
6	UK	70M	25T
7	USA	20M	30T
8	China	70M	25T
9	France	50M	3T
10	China	70M	25T

```
In [3]: salary_df = ps.DataFrame(
        {
            "Department": ["Finance","Technology","Finance","Technology","Technology"],
            "Staff": ["Tom", "Peter","Simon", "Mary", "Jane"],
            "Salary": [90000.00, 57000.00,40000.00, 34000.00, 12000.00]
        },
    )

salary_df
```

Out[3]:

	Department	Staff	Salary
0	Finance	Tom	90000.0
1	Technology	Peter	57000.0
2	Finance	Simon	40000.0
3	Technology	Mary	34000.0
4	Technology	Jane	12000.0

Create Spark DataFrame

```
In [4]: spark = SparkSession.builder.getOrCreate()
sdf=spark.createDataFrame([['France','50M','3T'], ['India','30M','30T'], ['Kenya','70M','25T'],
                           ['Nigeria','90M','60T'], ['China','20M','2T'], ['USA','80M','30T'],
                           ['UK','70M','25T'], ['USA','20M','30T'], ['China','70M','25T'],
                           ['France','50M','3T'], ['China','70M','25T'] ],
                           schema='Country string,Population string,GDP string')

sdf.show()
```

```
+-----+-----+---+
|Country|Population|GDP|
+-----+-----+---+
| France|      50M| 3T|
|  India|      30M|30T|
|  Kenya|     70M|25T|
|Nigeria|     90M|60T|
|   China|      20M| 2T|
|     USA|      80M|30T|
|     UK|      70M|25T|
|     USA|      20M|30T|
|   China|      70M|25T|
| France|      50M| 3T|
|  China|      70M|25T|
+-----+-----+---+
```

Create Pandas DataFrame

```
In [5]: pd_df=pd.DataFrame([[ 'France', '50M', '3T'],[ 'India', '30M', '30T'],[ 'Kenya', '70M', '25T'],
                             [ 'Nigeria', '90M', '60T'],[ 'China', '20M', '2T'],[ 'USA', '80M', '30T'],
                             [ 'UK', '70M', '25T'],[ 'USA', '20M', '30T'],[ 'China', '70M', '25T'],
                             [ 'France', '50M', '3T'],[ 'China', '70M', '25T'] ],
                             columns=[ 'Country', 'Population', 'GDP'])

pd_df
```

Out[5]:

	Country	Population	GDP
0	France	50M	3T
1	India	30M	30T
2	Kenya	70M	25T
3	Nigeria	90M	60T
4	China	20M	2T
5	USA	80M	30T
6	UK	70M	25T
7	USA	20M	30T
8	China	70M	25T
9	France	50M	3T
10	China	70M	25T

Read external csv file with Pandas-on-SPark

```
In [6]: df=ps.read_csv("titanic.csv")
```

```
In [7]: df.head()
```

Read csv data with Pandas-on-Saprk

```
In [8]: ps_to_pd_df=ps_df.to_pandas()  
ps_to_pd_df
```

Out[8]:

	Country	Population	GDP
0	France	50M	3T
1	India	30M	30T
2	Kenya	70M	25T
3	Nigeria	90M	60T
4	China	20M	2T
5	USA	80M	30T
6	UK	70M	25T
7	USA	20M	30T
8	China	70M	25T
9	France	50M	3T
10	China	70M	25T

Convert Pandas DataFrame to Pandas-on-Spark DataFrame

```
In [9]: pd_to_ps_df=ps.from_pandas(pd_df)
pd_to_ps_df
```

Out[9]:

	Country	Population	GDP
0	France	50M	3T
1	India	30M	30T
2	Kenya	70M	25T
3	Nigeria	90M	60T
4	China	20M	2T
5	USA	80M	30T
6	UK	70M	25T
7	USA	20M	30T
8	China	70M	25T
9	France	50M	3T
10	China	70M	25T

Convert Pandas DataFrame to Spark DataFrame

```
In [10]: pd_to_sdf=spark.createDataFrame(pd_df)
pd_to_sdf.show()
```

```
+-----+-----+---+
|Country|Population|GDP|
+-----+-----+---+
| France|      50M| 3T|
|  India|      30M|30T|
|  Kenya|     70M|25T|
|Nigeria|     90M|60T|
|   China|     20M| 2T|
|    USA|     80M|30T|
|    UK|     70M|25T|
|    USA|     20M|30T|
|   China|     70M|25T|
| France|     50M| 3T|
|  China|     70M|25T|
+-----+-----+---+
```

Pandas on Spark Functions

```
In [11]: salary_df
```

```
Out[11]:
```

	Department	Staff	Salary
0	Finance	Tom	90000.0
1	Technology	Peter	57000.0
2	Finance	Simon	40000.0
3	Technology	Mary	34000.0
4	Technology	Jane	12000.0

Check rows and columns


```
In [12]: salary_df.shape
```

```
Out[12]: (5, 3)
```

Check DataFrame types

```
In [13]: salary_df.info()
```

```
<class 'pyspark.pandas.frame.DataFrame'>
Int64Index: 5 entries, 0 to 4
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Department  5 non-null      object
1   Staff       5 non-null      object
2   Salary      5 non-null      float64
dtypes: float64(1), object(2)
```

Check Statistical Summary

```
In [14]: salary_df.describe()
```

Out[14]:

	Salary
count	5.000000
mean	46600.000000
std	29117.005341
min	12000.000000
25%	34000.000000
50%	40000.000000
75%	57000.000000
max	90000.000000

Calculate Sum

```
In [15]: salary_df['Salary'].sum()
```

Out[15]: 233000.0

Calculate Mean

```
In [16]: salary_df['Salary'].mean()
```

Out[16]: 46600.0

Calculate Standard Deviation

```
In [17]: salary_df['Salary'].std()
```

```
Out[17]: 29117.005340522228
```

Calculate Variance of Salary

```
In [18]: salary_df['Salary'].var()
```

```
Out[18]: 847800000.0
```

Calculate Skewnes of Salary

```
In [19]: salary_df['Salary'].skew()
```

```
Out[19]: 0.44342185901218767
```

Group Salary by Department

```
In [20]: salary_df.groupby('Department')['Salary'].sum()
```

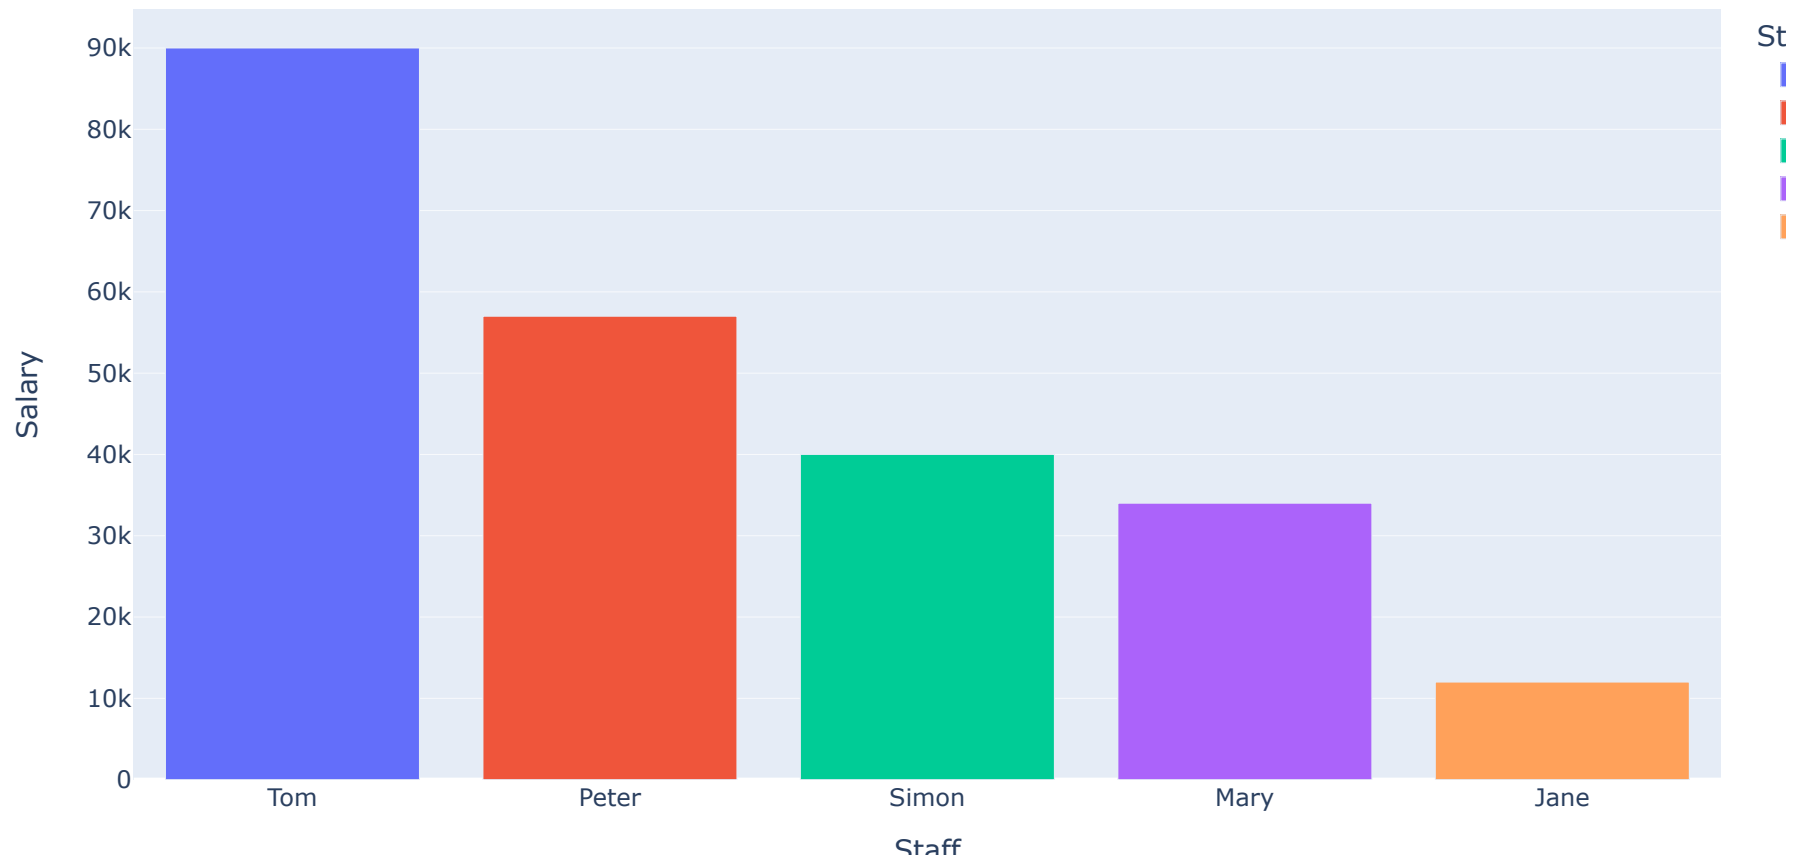
```
Out[20]: Department
Finance      130000.0
Technology   103000.0
Name: Salary, dtype: float64
```

Ploting Visualizations in Pandas on Spark

Pandas on Spark leverages plotly in the backend for visualization


Let's plot Salary for each Staff on a bar graph

```
In [21]: salary_df.plot.bar(x='Staff',y='Salary',color='Staff')
```



Pie Chart of the Salary per Department

```
In [22]: salary_df.groupby('Department')['Salary'].sum().plot.pie()
```



Kernel Density Estimation for a normal distribution data

```
In [23]: ps.DataFrame(np.random.normal(10,2,10000)).plot.kde(bw_method=3)
```

22/01/14 11:48:38 WARN InstanceBuilder\$NativeBLAS: Failed to load implementation from:dev.ludovic.netlib.blas.JNIBLAS
22/01/14 11:48:39 WARN InstanceBuilder\$NativeBLAS: Failed to load implementation from:dev.ludovic.netlib.blas.Foreign
LinkerBLAS

